Stability of the per-configuration difficulty ordering across training seeds

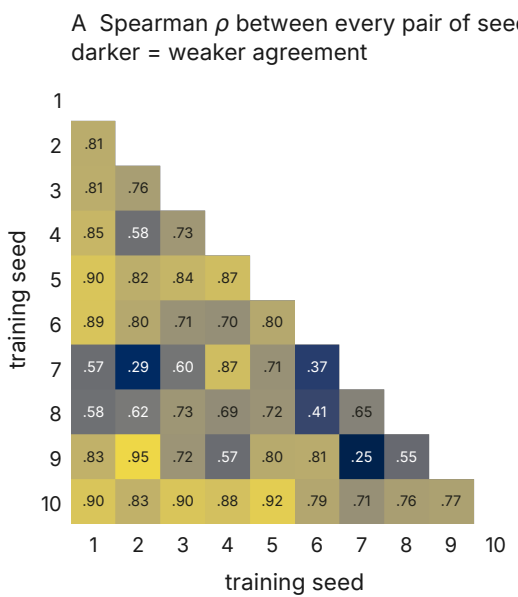

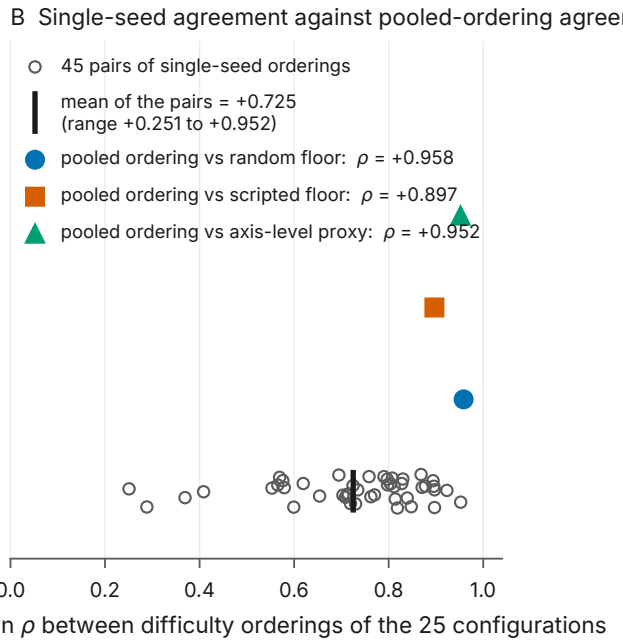

docs/gate1_result.json: per-seed, per-configuration mean damage taken, 200 evaluation episodes per cell. Panel B's pooled values use the seed-pooled ordering against docs/gate1_floors.json and the axis-level proxy defined in this script. The two quantities are both Spearman correlations over the same 25 configurations.